

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

DESCRIPTION

Background of Problem Statement :

The US Census Bureau has published California Census Data which has 10 types of metrics such as the population, median income, median housing price, and so on for each block group in California. The dataset also serves as an input for project scoping and tries to specify the functional and nonfunctional requirements for it.

Problem Objective :

The project aims at building a model of housing prices to predict median house values in California using the provided dataset. This model should learn from the data and be able to predict the median housing price in any district, given all the other metrics.

Districts or block groups are the smallest geographical units for which the US Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people). There are 20,640 districts in the project dataset.

```
[40]: #Import Necessary Libraries:
import pandas as pd
import numpy as np

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.tree import DecisionTreeRegressor
import statsmodels.formula.api as smf

from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from matplotlib.axes._axes import _log as matplotlib_logger
matplotlib_logger.setLevel('ERROR')
```

Simple 0 1 Python 3 Idle

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

Build a model of housing prices to predict median house values in California using the provided dataset.

Load the data :

Read the 'housing.csv' file from the folder into the program.

Print first few rows of this data.

```
[41]: df_house = pd.read_excel('1553768847_housing.xlsx')
[42]: df_house.head()
[43]: longitude latitude housing_median_age total_rooms total_bedrooms population households median_income ocean_proximity median_house_value
0 -122.23 37.88 41 880 128.0 322 126 8.3252 NEAR BAY 452600
1 -122.22 37.86 21 7099 1106.0 2401 1138 8.3014 NEAR BAY 358500
2 -122.24 37.85 52 1467 190.0 496 177 7.2574 NEAR BAY 352100
3 -122.25 37.85 52 1274 235.0 558 219 5.6431 NEAR BAY 341300
4 -122.25 37.85 52 1627 280.0 565 259 3.6462 NEAR BAY 342200
[44]: import math
print(math.log(452600))
13.022764812181574
[45]: df_house.columns
[46]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
'total_bedrooms', 'population', 'households', 'median_income',
'ocean_proximity', 'median_house_value'],
dtype='object')
Handle missing values :
```

Simple 0 1 Python 3 Idle

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

Handle missing values :

Fill the missing values with the mean of the respective column.

```
[45]: df_house.isnull().sum()
[46]: longitude 0
latitude 0
housing_median_age 0
total_rooms 0
total_bedrooms 267
population 0
households 0
median_income 0
ocean_proximity 0
median_house_value 0
dtype: int64
[47]: df_house.total_bedrooms = df_house.total_bedrooms.fillna(df_house.total_bedrooms.mean())
df_house.isnull().sum()
[48]: longitude 0
latitude 0
housing_median_age 0
total_rooms 0
total_bedrooms 0
population 0
households 0
median_income 0
ocean_proximity 0
median_house_value 0
dtype: int64
Encode categorical data :
Convert categorical column in the dataset to numerical data.
```

Simple 0 1 Python 3 Idle

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

Convert categorical column in the dataset to numerical data.

```
[47]: le = LabelEncoder()
df_house['ocean_proximity'] = le.fit_transform(df_house['ocean_proximity'])
```

Standardize data:

Standardize training and test datasets.

```
[52]: names = df_house.columns

scaler = StandardScaler()
scaled_df = scaler.fit_transform(df_house)
scaled_df = pd.DataFrame(scaled_df, columns=names)
scaled_df.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	median_house_value
0	-1.327835	1.052548	0.982143	-0.804819	-0.975228	-0.974429	-0.977033	2.344766	1.291089	2.129631
1	-1.322844	1.043185	-0.607019	2.045890	1.335088	0.861439	1.669961	2.332238	1.291089	1.314156
2	-1.332827	1.038503	1.856182	-0.535746	-0.829732	-0.820777	-0.843637	1.782699	1.291089	1.256693
3	-1.337818	1.038503	1.856182	-0.624215	-0.722399	-0.766028	-0.733781	0.932968	1.291089	1.165100
4	-1.337818	1.038503	1.856182	-0.462404	-0.615086	-0.759847	-0.629157	-0.012881	1.291089	1.172900

Extract input (X) and output (Y) data from the dataset.

```
[53]: X_features = ['longitude', 'latitude', 'housing_median_age', 'total_rooms',
                'total_bedrooms', 'population', 'households', 'median_income',
                'ocean_proximity']
X_scaled_df[X_features]
Y_scaled_df['median_house_value']
```

Simple 0 1 Python 3 Idle

Mode: Command Ln 1, Col 1 California Housing Price Prediction.py:nb

Used 4.5 of 50 hours in May, 2023

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

```
Y_scaled_df['median_house_value']

print(type(X))
print(type(Y))

<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

Split the dataset:

```
[54]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=1)

print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)

(16512, 9) (16512,)
(4128, 9) (4128,)
```

Perform Linear Regression:

Perform Linear Regression on training data.

Predict output for test dataset using the fitted model.

Print root mean squared error (RMSE) from Linear Regression. [HINT: Import mean_squared_error from sklearn.metrics]

```
[55]: linreg = LinearRegression()
linreg.fit(X_train, y_train)

[55]: LinearRegression()

[56]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

[56]: LinearRegression(normalize=False)

[57]: y_predict = linreg.predict(X_test)
```

Simple 0 1 Python 3 Idle

Mode: Command Ln 1, Col 1 California Housing Price Prediction.py:nb

Used 4.5 of 50 hours in May, 2023

Used 4.5 of 50 hours in May, 2023

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

```
[58]: print(sqrt(mean_squared_error(y_test, y_predict)))
print((r2_score(y_test, y_predict)))

0.6856598120301221
0.6276223517959295
```

Perform Decision Tree Regression:

Perform Decision Tree Regression on training data.

Predict output for test dataset using the fitted model.

Print root mean squared error from Decision Tree Regression.

```
[59]: dtreg = DecisionTreeRegressor()
dtreg.fit(X_train, y_train)

[59]: DecisionTreeRegressor()
```

```
[61]: y_predict = dtreg.predict(X_test)
print(sqrt(mean_squared_error(y_test, y_predict)))
print((r2_score(y_test, y_predict)))

0.59673596861619542
0.6385147998268702
```

Perform Random Forest Regression:

Perform Random Forest Regression on training data.

Predict output for test dataset using the fitted model.

Print RMSE (root mean squared error) from Random Forest Regression.

```
[63]: from sklearn.ensemble import RandomForestRegressor
rfreg = RandomForestRegressor()
```

Simple 0 1 Python 3 Idle

Mode: Command Ln 1, Col 1 California Housing Price Prediction.py:nb

Used 4.5 of 50 hours in May, 2023

Jupyter

Used 4.5 of 50 hours in May, 2023

Start Lab End Lab

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

Filter files by name

Name	Last Modified
155376884...	21 hours ago
California H...	12 minutes ago

```

[63]: RandomForestRegressor()

[66]: rfreg=RandomForestRegressor()
      rfreg.fit(x_train,y_train)

[66]: RandomForestRegressor()

[69]: y_predict = rfreg.predict(x_test)
      print(sort(mean_squared_error(y_test,y_predict)))
      print((r2_score(y_test,y_predict)))
      0.4256352518189598
      0.8168917256841964

Perform Linear Regression with one independent variable :

Extract just the median_income column from the independent variables (from X_train and X_test).

Perform Linear Regression to predict housing values based on median_income.

Predict output for test dataset using the fitted model.

Plot the fitted model for training data as well as for test data to check if the fitted model satisfies the test data.

[70]: X_train_Income=x_train[['median_income']]
      X_test_Income=x_test[['median_income']]

[71]: print(x_train_Income.shape)
      print(y_train.shape)
      (16512, 1)
      (16512,)

[72]: linreg=LinearRegression()
      linreg.fit(x_train_Income,y_train)
      y_predict = linreg.predict(x_test_Income)

```

Simple 0 1 Python 3 Idle

Mode: Command Ln 1, Col 1 California Housing Price Prediction.py

Jupyter

Used 4.5 of 50 hours in May, 2023

Start Lab End Lab

File Edit View Run Kernel Tabs Settings Help

California Housing Price Prediction

Filter files by name

Name	Last Modified
155376884...	21 hours ago
California H...	12 minutes ago

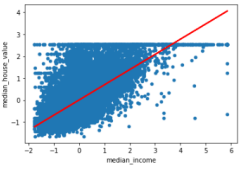
```

[73]: print(linreg.intercept_, linreg.coef_)
      print(sort(mean_squared_error(y_test,y_predict)))
      print((r2_score(y_test,y_predict)))
      0.005623019866893164 [0.69338221]
      0.7212595914243148
      0.47198835934467734

[74]: scaled_df.plot(kind='scatter',x='median_income',y='median_house_value')
      plt.plot(x_test_Income,y_predict,c='red',linewidth=2)

[74]: [matplotlib.lines.Line2D at 0x746e10b290b]

```



The figure is a scatter plot showing the relationship between median_income (x-axis) and median_house_value (y-axis). The x-axis ranges from -2 to 6, and the y-axis ranges from -1 to 4. The plot contains numerous blue data points representing individual housing units. A solid red line is overlaid on the scatter plot, representing the linear regression model. The red line shows a positive correlation, indicating that as median income increases, the median house value also tends to increase. The data points are most densely clustered between median_income values of 1 and 3, and median_house_value values of 1 and 2.

Simple 0 1 Python 3 Idle

Mode: Command Ln 1, Col 1 California Housing Price Prediction.py